

(How do we find the subtypes of a disease? –)

Clustering and designing new k-means  
algorithm variants via probabilistic modeling

Optional !Saturday Track!  
CS5691 PRML Jul-Nov 2025

(used previously in “Algorithmic Thinking in Bioinformatics”  
Course

Manikandan Narayanan  
Associate Professor, Computer Science and Engineering, IIT  
Madras)

# Context: Seen so far...

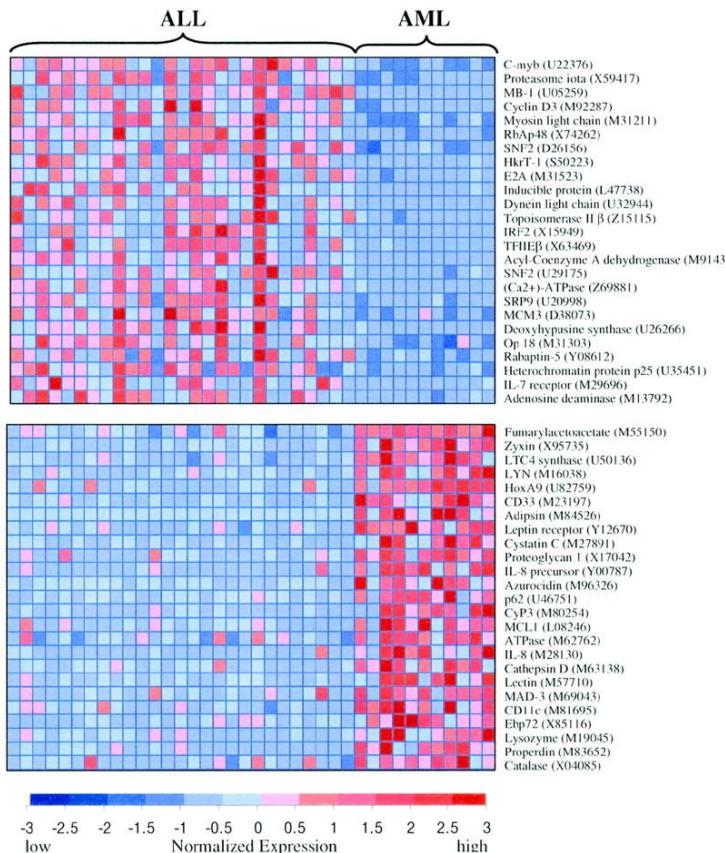
- Session on Motivation and Background for “Algorithmic Thinking in Bioinformatics” course
- Sessions on several biological questions & their algorithmic solutions
- This session: Think about a
  - clinical question (“how do we identify the different (sub)types of a disease?”),
  - associated computational problem (clustering), and
  - its algorithmic solution(s) (k-means algorithm and designing this algo.’s new variants using probabilistic modeling).

# Cancer example: Can we group cancer samples into different (sub)types?



Golub et al. Science 1999

# Cancer example: from grouping to prediction!



Golub et al. Science 1999

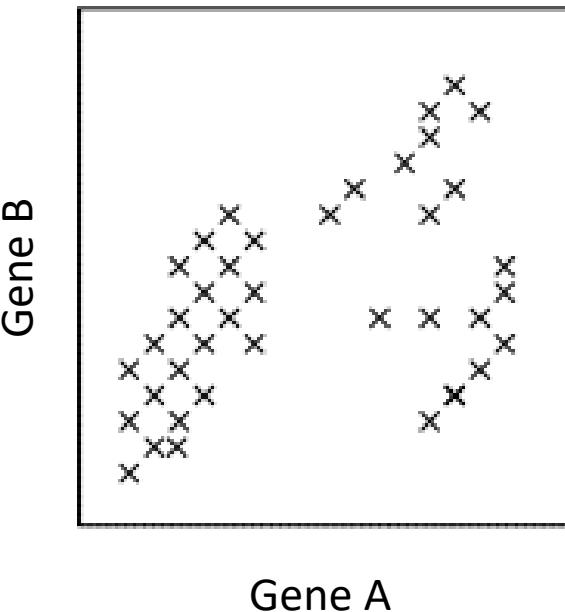
# Clustering: a popular data science task

- Human brains are good at grouping objects based on their similarity. How do we automate it?
  - Group N objects into K groups
  - Paradigms:
    - K-means clustering – centroid-based;
    - Hierarchical clustering – nested/stratified with top-down or bottom-up approaches;
    - Spectral clustering – graph connectivity-based;
    - etc.
- Why is a good clustering popular? It
  - allows exploratory data analysis (unsupervised machine learning)
  - has predictive power
  - allows lossy compression
  - can reveal interesting outliers

# Clustering: a popular data science task viewed thru' a probabilistic modeling lens

- Our approach in this lecture
  - Present k-means approach to clustering, starting with a non-probabilistic approach (hard k-means) and then transitioning to a probabilistic approach (soft k-means).
  - *Focus is on viewing (k-means) clustering as an inference+parameter-learning task based on a probabilistic mixture model – this mixture dens. estn. approach will help us design new k-means algo. variants!*
- Sources for this lecture:
  - Main source:
    - David J. C. MacKay. Information Theory, Inference and Learning Algorithms (Chapters 20-22). 2003. – **cited as [DJM]** for content/figures taken as is from this book.
  - Other sources:
    - Christopher M. Bishop. Pattern Recognition and Machine Learning. 2006. – **cited as [CMB]** for figures taken as is from this book.
    - Stanford CS228 Probabilistic Graphical Models Course Lecture Notes.  
<https://ermongroup.github.io/cs228-notes/> - cited as **[ECL]**

$N=40$  data points over two dimensions (genes)



# K-means Clustering Outline

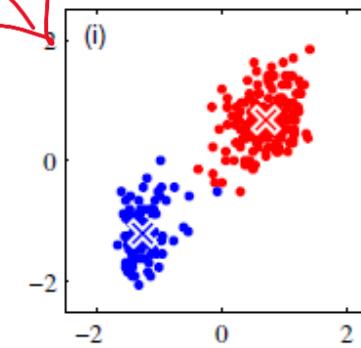
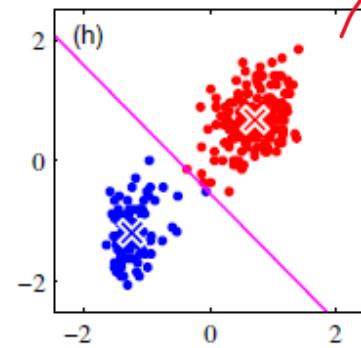
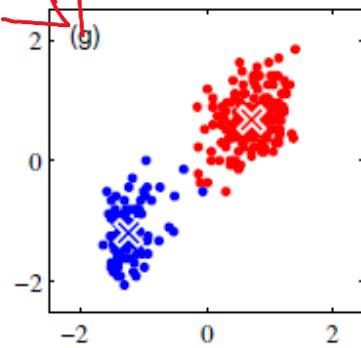
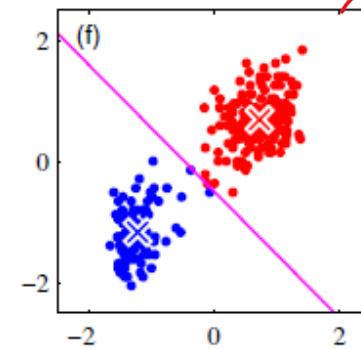
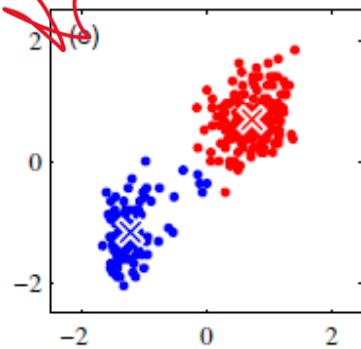
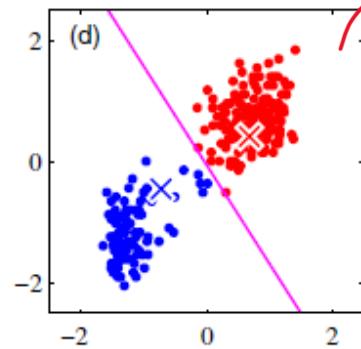
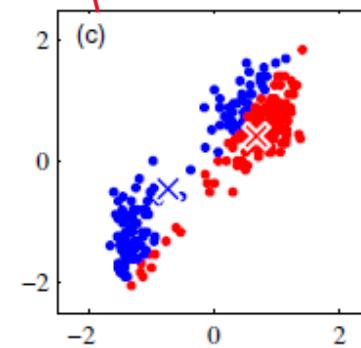
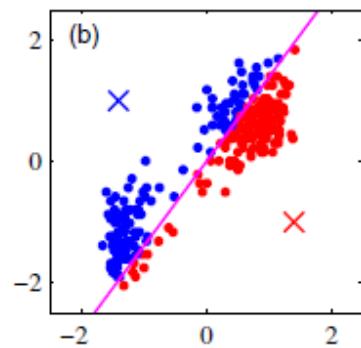
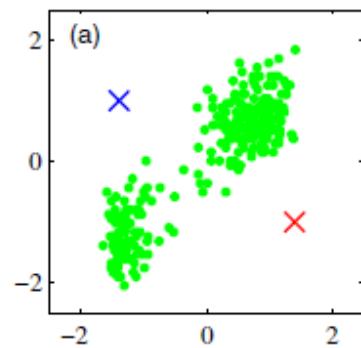
- 1) hard version: algorithm and analysis**
- 2) soft version vanilla (aided by intuitive formula)
- 3) soft version vanilla (aided by a vanilla probabilistic model)
- 4) soft version enhancements (aided by an enhanced probabilistic model)

# K-means clustering (in one slide!)

Problem: Find K cluster centers that minimize the sum of squared distance of each data point to the nearest cluster center.

Algorithm (Lloyd's Heuristic): Starting with K centers (means) initialized in some way, iterate until convergence:

- a) **[Centers -> Clusters]** *Assign* each data point to the nearest center.
- b) **[Clusters -> Centers]** *Update* centers to sample means of data points they are responsible for.



# K-means clustering problem and notation

**Input:**  $N$  objects:  $\{x^{(n)}\} = x^{(1)}, x^{(2)}, \dots, x^{(N)}$

**Output:**  $K$  centers/means:  $\{m^{(k)}\} = m^{(1)}, m^{(2)}, \dots, m^{(K)}$

$$m^{(k)} \in \mathbb{R}^I$$

**(Squared) Distance measure:**  $d(x, m) = \sum_{i=1}^I (x_i - m_i)^2$ , where  $I$  is dimensionality of data (2 in our examples).

**Problem:** Find  $K$  means that minimize  $\sum_{n=1}^N d(x^{(n)}, m^{(k^*(n))})$ , where  $k^*(n) = \operatorname{argmin}_k d(x^{(n)}, m^{(k)})$

# K-means clust. pseudocode: Lloyd's heuristic

**Initialization:** Set  $K$  means  $\{m^{(k)}\}$  to random values.

**Assignment:**  $r_k^{(n)} = \begin{cases} 1 & \text{if } m^{(k)} \text{ is the closest mean to datapoint } x^{(n)} \\ 0 & \text{otherwise} \end{cases}$

**Update:**

$$m^{(k)} = \frac{\sum_n r_k^{(n)} x^{(n)}}{R^{(k)}}.$$

where  $R^{(k)}$  is the total responsibility of mean  $k$ ,

$$R^{(k)} = \sum_n r_k^{(n)}.$$

# K-means clust. pseudocode: Lloyd's heuristic

**Initialization:** Set  $K$  means  $\{m^{(k)}\}$  to random values.

**Assignment:**  $r_k^{(n)} = 1$       if  $m^{(k)}$  is the closest mean to datapoint  $x^{(n)}$   
                  0      otherwise

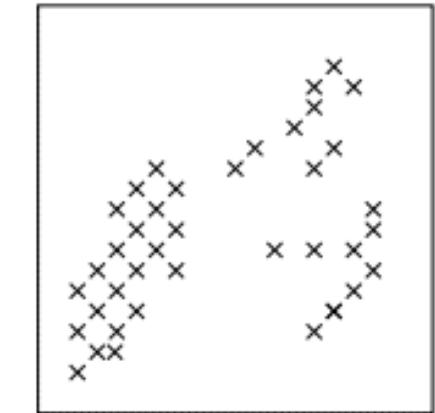
**Update:**

$$m^{(k)} = \frac{\sum_n r_k^{(n)} x^{(n)}}{R^{(k)}} .$$

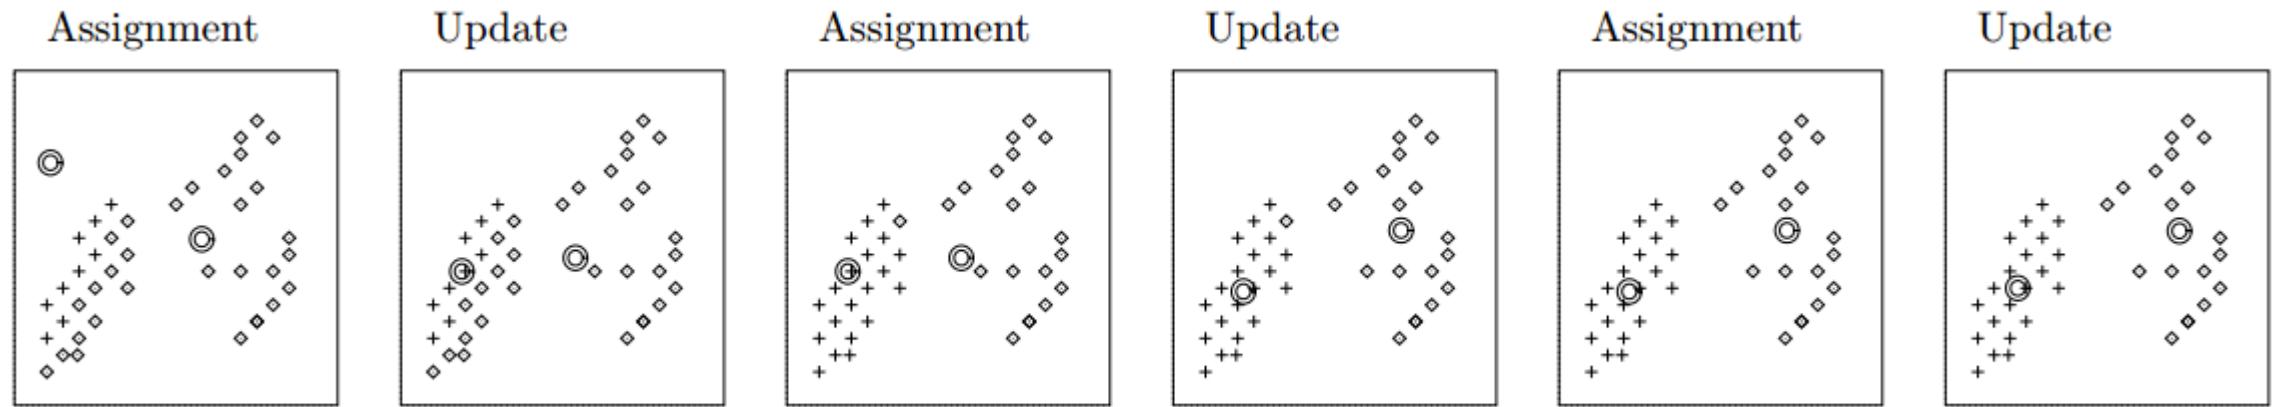
where  $R^{(k)}$  is the total responsibility of mean  $k$ ,

$$R^{(k)} = \sum_n r_k^{(n)} . \quad \sum_k R^{(k)} = N$$

# Example run ( $K=2$ )

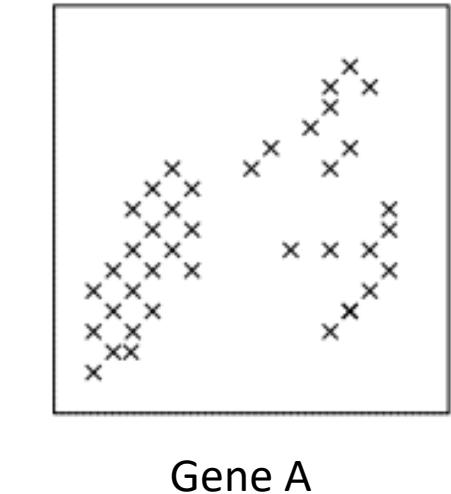


Gene A

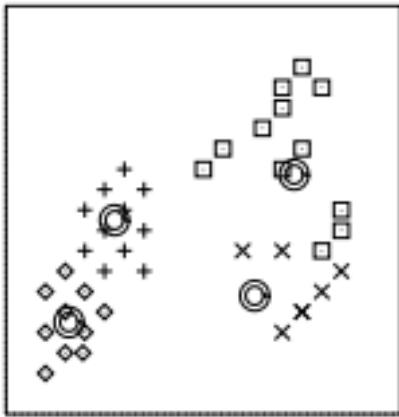
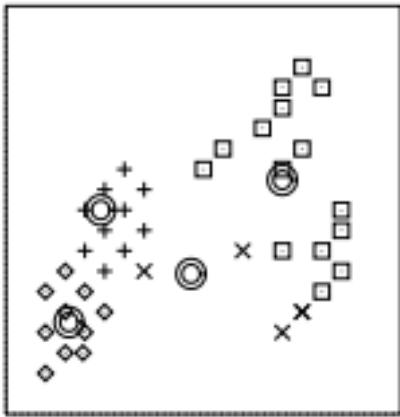
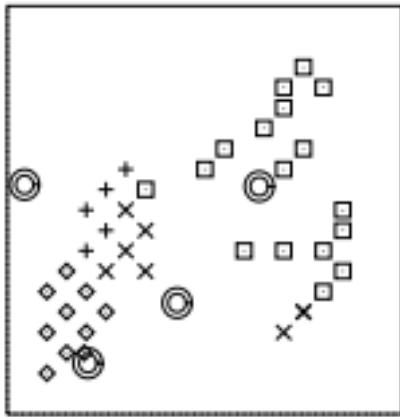


Let's try K=4 clusters  
with two different starting points!

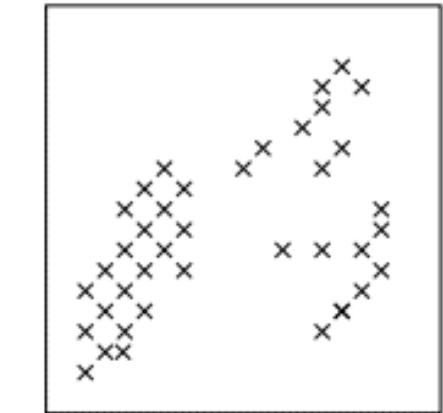
# Example run1 ( $K=4$ )



Run 1

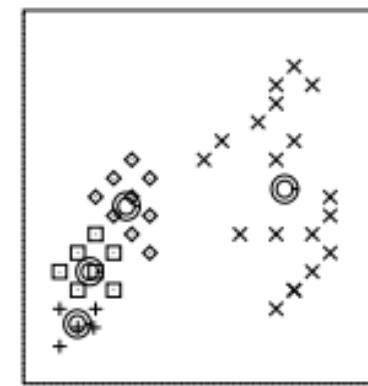
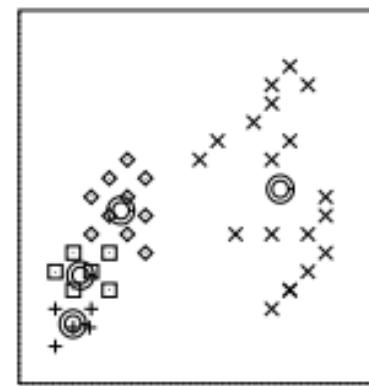
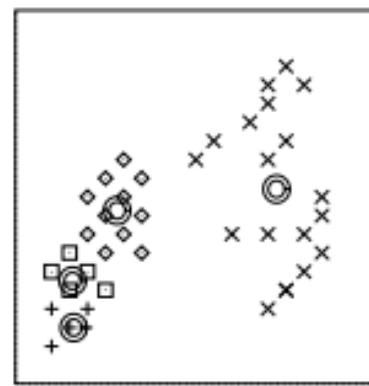
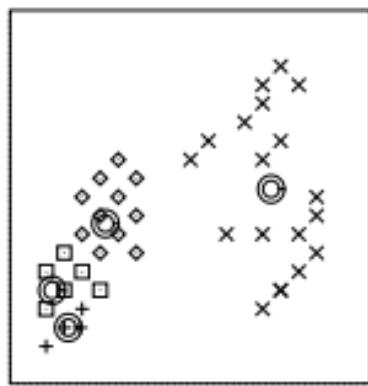
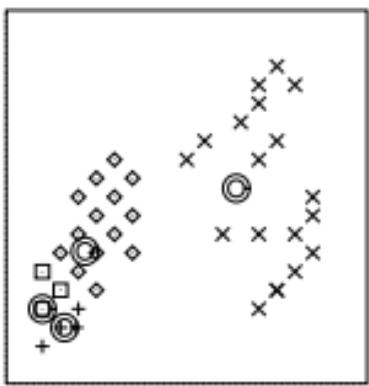
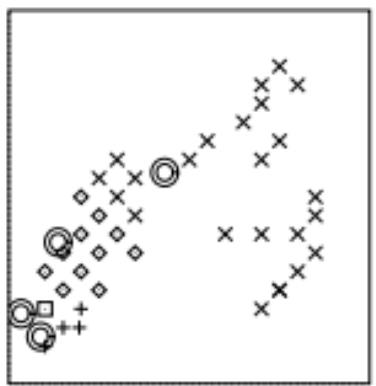


# Example run2 ( $K=4$ )



Gene A

Run 2



# Analysis of algorithm – Questions

**Problem:** Find K cluster centers that minimize the sum of squared distance of each data point to the nearest cluster center.

- 1) Are the means (centroids) the best (and **only**) cluster centers for a given partition?

$$\text{Is } m^{(k)} = \arg \min_p \sum_{n: r_n^{(k)}=l} d(x^{(n)}, p) \text{ ?}$$

- 2) Does the algorithm converge always for any dataset?

- 3) Does the algorithm actually (globally) minimize the sum of squared distances?

# K-means clust. pseudocode: Lloyd's heuristic

**Initialization:** Set  $K$  means  $\{m^{(k)}\}$  to random values.

**Assignment:**  $r_k^{(n)} = 1$  if  $m^{(k)}$  is the closest mean to datapoint  $x^{(n)}$   
0 otherwise

$$\text{objfn} = \sum_n \sum_k r_k^{(n)} d(x^{(n)}, m^{(k)})$$

**Update:**

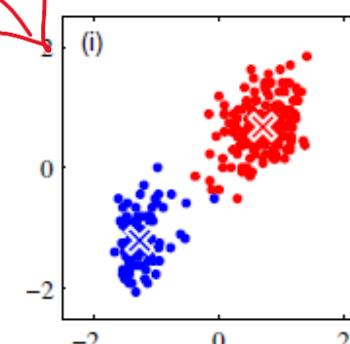
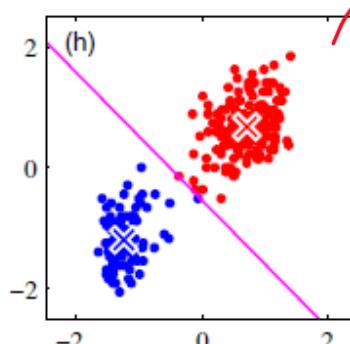
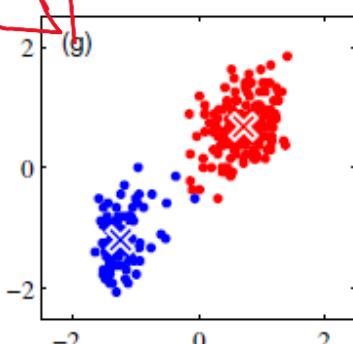
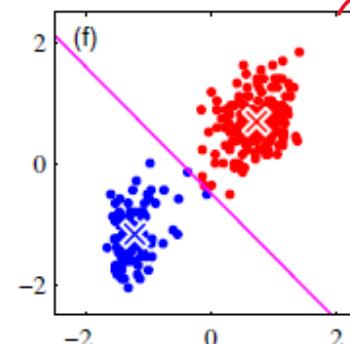
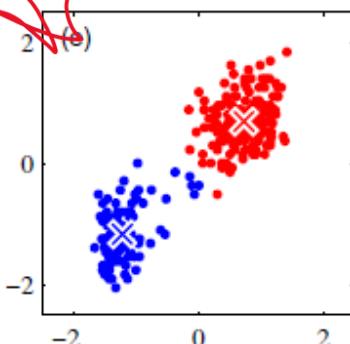
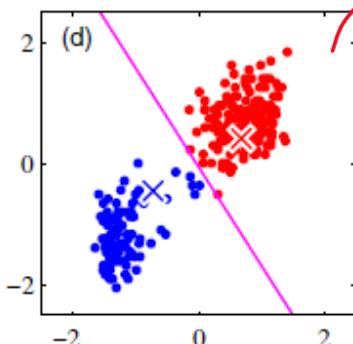
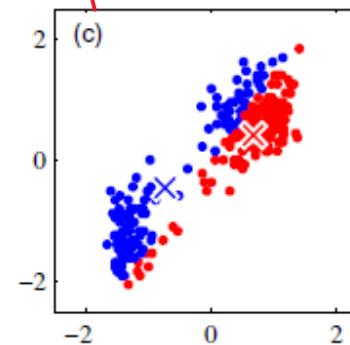
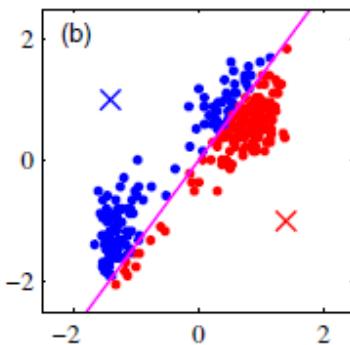
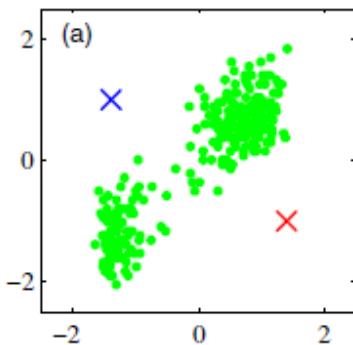
$$m^{(k)} = \frac{\sum_n r_k^{(n)} x^{(n)}}{R^{(k)}}$$

(centroid or center  
of gravity of cluster k)  
is the only minimizer

where  $R^{(k)}$  is the total responsibility of mean  $k$ ,

$$R^{(k)} = \sum_n r_k^{(n)}.$$

obj fn decrease



[CMB]

# Analysis of algorithm – Answers

**Problem:** Find K cluster centers that minimize the sum of squared distance of each data point to the nearest cluster center.

1) Are the means (centroids) the best (and only) cluster centers for a given partition?

Yes, cluster k centroid is the only minimizer of cluster k cost, given a partition. (Prove it!)

2) Does the algorithm converge always for any dataset?

Yes, because:  
a) Every iteration reduces  $\text{obj. fn}$ .  
b)  $\text{obj. fn}$  is lower-bounded.  
c) There are only finite # of ways to partition N datapts into K clusters.

3) Does the algorithm actually (globally) minimize the sum of squared distances?

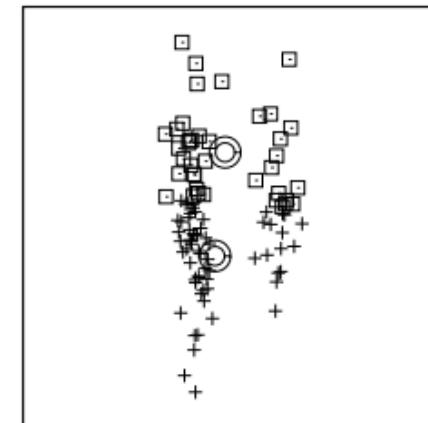
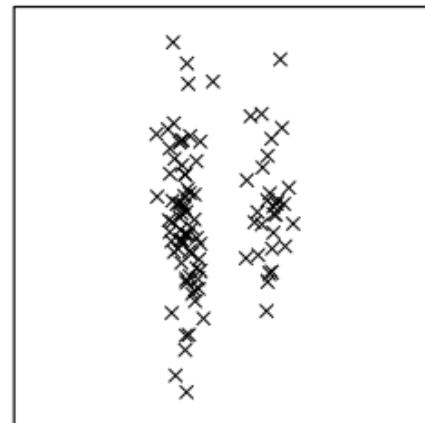
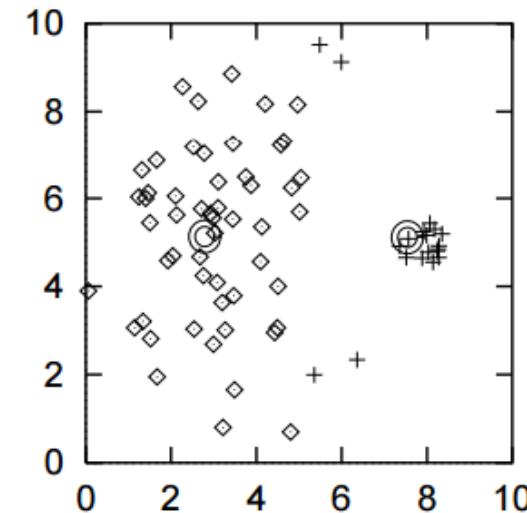
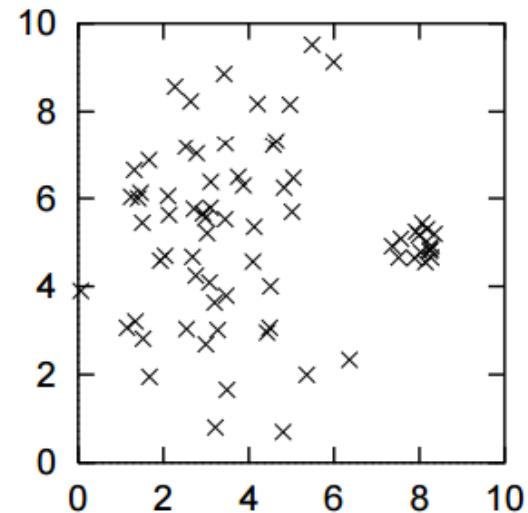
No, only reaches local minima as seen in k=4 cluster runs above. (NP-hard to find global minima)

# Food for thought – Exercise questions

- Prove that (hard) k-means yields clusters that are each convex sets!
- Prove that k-means is efficiently (polynomial-time) solvable when all datapoints lie on a real line (i.e., finding k centers that minimize k-means cost function when all datapoints are 1D).

Blank space

# K-means pitfalls



# K-means pitfalls: solutions?

- How about points in the border between two means? THINK FUZZY.
- What is a better alternative for distance  $d(\cdot, \cdot)$ ? THINK ELLIPSOIDAL.
  - Instead of giving equal weight to all dimensions in a spherical fashion as in squared Euclidean distance, why not choose a distance that adapts to the width and breadth of each cluster?

# K-means Clustering Outline

- 1) hard version: algorithm and analysis
- 2) soft version vanilla (aided by intuitive formula)**
- 3) soft version vanilla (aided by a vanilla probabilistic model)
- 4) soft version enhancements (aided by an enhanced probabilistic model)

Q: Can we make  $r_k^{(n)}$  fuzzy/soft using an intuitive formula?

Q: Can we make  $r_k^{(n)}$  fuzzy/soft using an intuitive formula?

- Intuitive formula (non-probabilistic or heuristic methods):

$$\frac{d_k}{d_1 + d_2 + d_3} \xrightarrow{\text{(or)}} \frac{e^{-\beta d_k}}{e^{-\beta d_1} + e^{-\beta d_2} + e^{-\beta d_3}}$$

# Vanilla Soft K-means Clustering - pseudocode

**Initialization:** Set  $K$  means  $\{m^{(k)}\}$  to random values.

**Assignment:** (E-step)  $r_k^{(n)} = \frac{\exp(-\beta d(\mathbf{m}^{(k)}, \mathbf{x}^{(n)}))}{\sum_{k'} \exp(-\beta d(\mathbf{m}^{(k')}, \mathbf{x}^{(n)}))}.$

**Update:** (M-step)  $\mathbf{m}^{(k)} = \frac{\sum_n r_k^{(n)} \mathbf{x}^{(n)}}{R^{(k)}}$

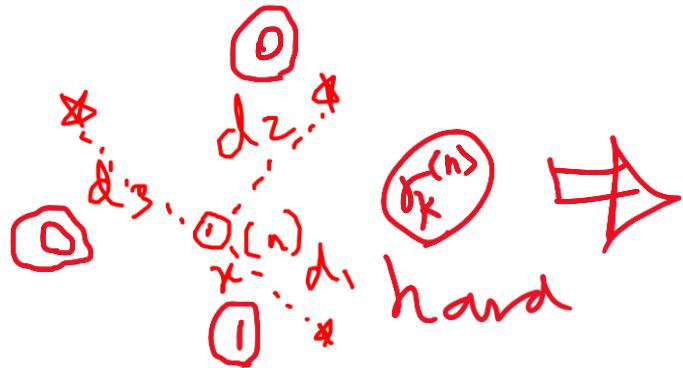
where  $R^{(k)}$  is the total responsibility of mean  $k$ ,

$$R^{(k)} = \sum_n r_k^{(n)}.$$

Blank space for illustrations: What is the role of the *stiffness/hardness* parameter  $\beta$ ?

Does this formula have any connection to a probabilistic model, and can it be exploited?

- Intuitive formula:



$$\frac{e^{-\beta d_k}}{e^{-\beta d_1} + e^{-\beta d_2} + e^{-\beta d_3}}$$

soft

$$p_k^{(n)} = P(z^{(n)}=k \mid X=x^{(n)}, \theta)$$

- Yes, this formula corresponds to an inference calculation for a probabilistic mixture model.
- Yes, understanding this connection can result in systematic design of new k-means algo. variants!

# K-means Clustering Outline

- 1) hard version: algorithm and analysis
- 2) soft version vanilla (aided by intuitive formula)
- 3) soft version vanilla (aided by a vanilla probabilistic model)**
  - 3a) (Brief) Background: Mixture density estimation via EM algorithm**
  - 3b) Methodology: Soft k-means heuristic/algorithm
- 4) soft version enhancements (aided by an enhanced probabilistic model)

# Recall: Example mixture density (very brief mention)

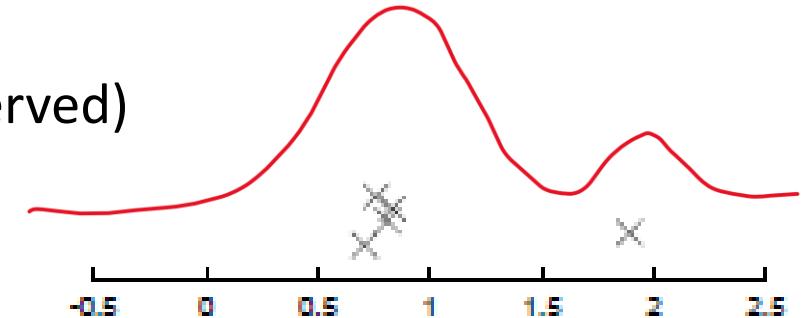
Mixture Model (MM) or Latent Variable Model (LVM)

$$Z \sim \text{Bernoulli}(\pi)$$

(latent variable, not observed)

$$(X | Z = z) \sim \mathcal{N}(\mu_z, \sigma_z^2)$$

( $X$  is observed)



MLE of  $\theta = (\boldsymbol{\pi}, \mu_1, \sigma_1, \mu_2, \sigma_2)$  based on:

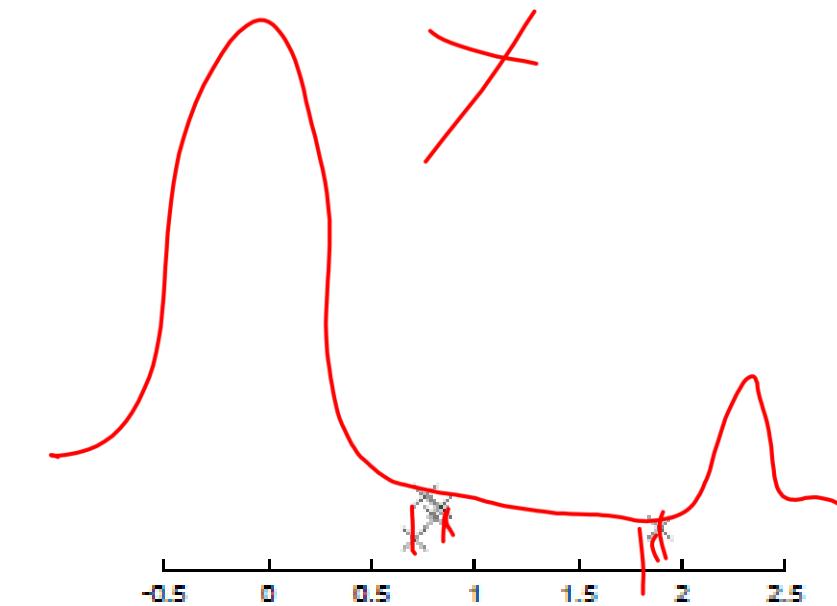
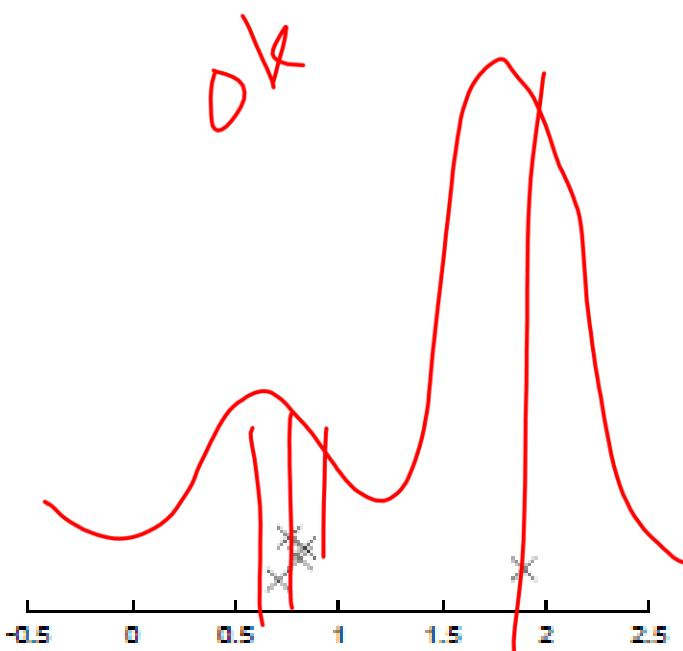
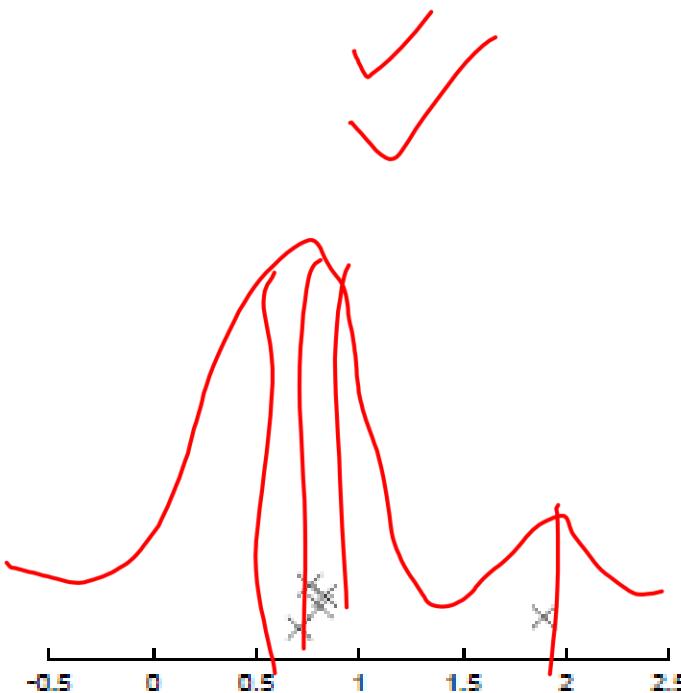
$$\mathcal{L}(\theta | D_N) = \prod_{n=1}^N (\boldsymbol{\pi} \mathcal{N}(x_n | \mu_1, \sigma_1^2) + (1 - \boldsymbol{\pi}) \mathcal{N}(x_n | \mu_2, \sigma_2^2))$$

$$\hat{\theta}_{MLE} = \arg \max_{\theta} \mathcal{L}(\theta | D_N) \quad (\text{using numerical methods like Newton-Raphson, or Expectation Maximization (EM) algorithm})$$

# What is optimized in the MLE of a mixture model?

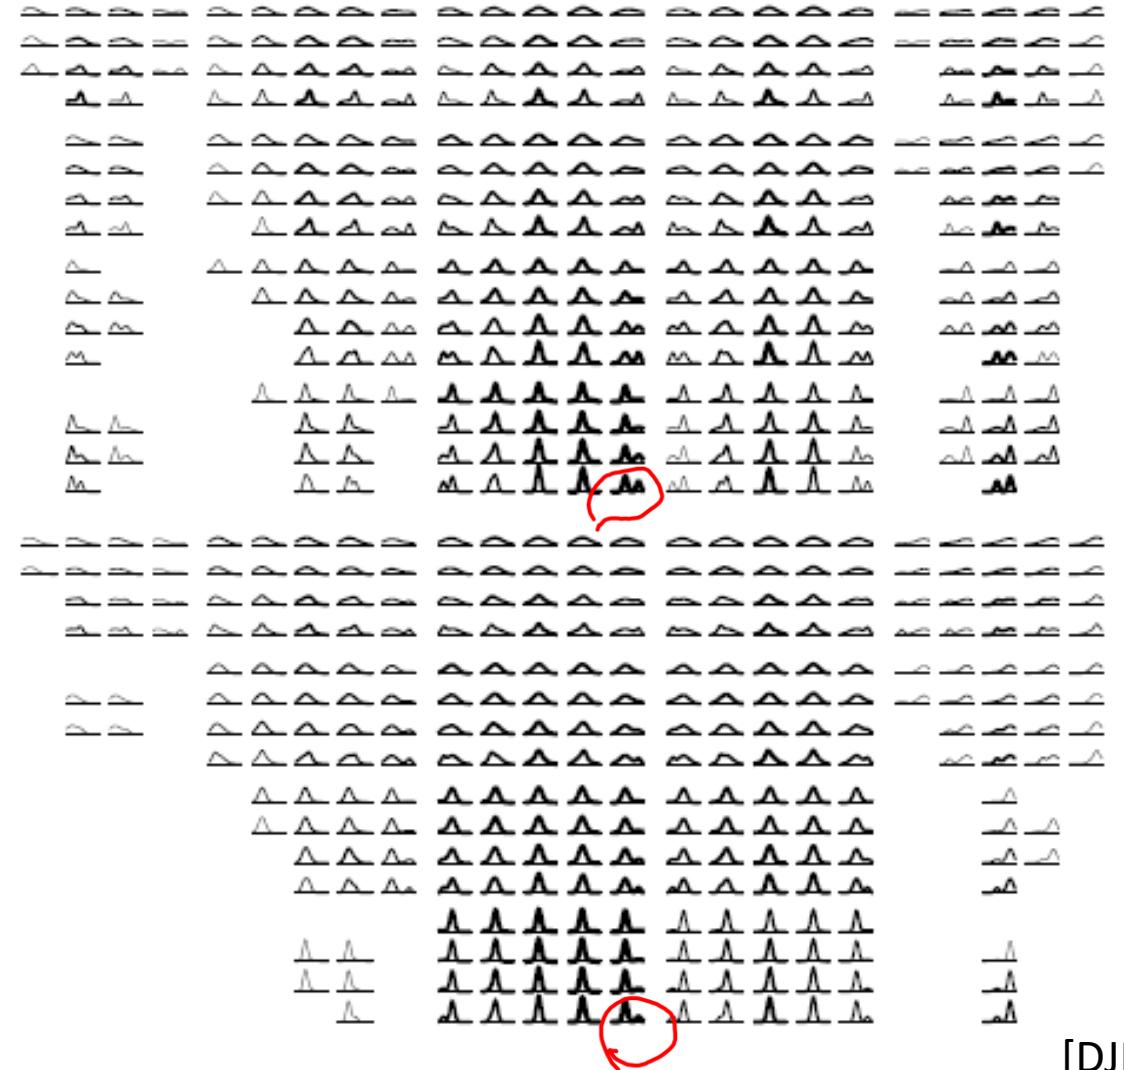
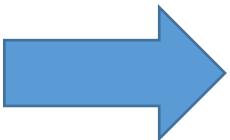
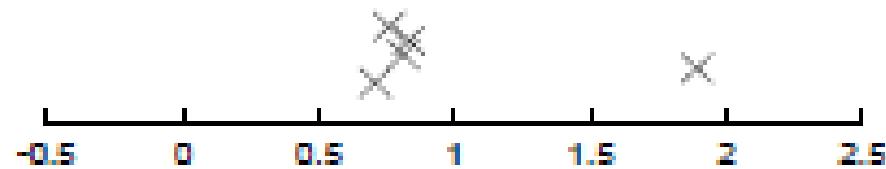
$$\boldsymbol{\theta} = (\mu_1, \sigma_1^2, \mu_2, \sigma_2^2, \pi_1)$$

5 params.



# Goal: MLE for a mixture of two 1D Gaussians ("Visual" soln.)

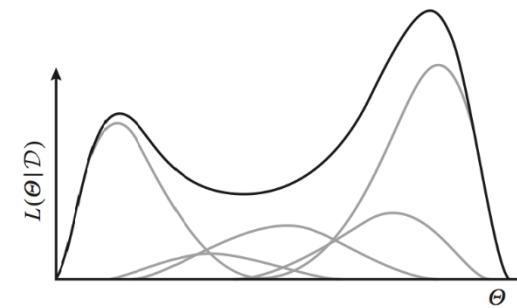
Estimate 5 params. now compared to  
2 parameters before!



[DJM]

# From visual to automated learning of a mixture model: Why is it hard?

$$\begin{aligned}\hat{\theta}_{ML} &= \arg \max_{\theta} \log (P(\mathbf{x}; \theta)) \\ &= \arg \max_{\theta} \log (\sum_{\mathbf{z}} P(\mathbf{x}, \mathbf{z}; \theta))\end{aligned}$$



- “log of sum”
  - necessitates an approximate learning approach\*
    - \*(fineprint: though GMMs are a tractable special case studied by theoretical CS community as well, we use them to motivate a generic algo. for mixture model density estimation.)
  - EM algo. is one such approach that pushes the log across the sum!
    - E-step: do inference
    - M-step: do estimation of params. using the inferred probabs.

Example (for GMM vs. 1D Gaussian): Why is it hard to optimize using basic calculus/optn.?

The diagram illustrates the optimization of two different models. On the left, under the heading "GMM: coupled eqns", there is a box containing the function  $f: LL(\theta; D)$  and several partial derivative equations:  $\frac{\partial f}{\partial \mu_1} = 0$ ,  $\frac{\partial f}{\partial \sigma^2} = 0$ , and  $\dots$ . A handwritten note below the box says "GMM: coupled eqns". On the right, under the heading "1D Gauss", there is a box containing the function  $f: f(\sigma) = 0$  and the derivative  $\frac{df}{d\sigma} = 0$ . Below these, there is a formula for the maximum likelihood estimate:  $\hat{\mu}_{ML} = \frac{\sum x_n}{N}$ . To the right of this formula is another equation:  $g(\hat{\sigma}, \hat{\mu}) = 0$ .

A **Newton-Raphson** algorithm for function optimization is possible for GMM estimation -- see Exercise 22.5 (Pg. 302 of [DJM]) – but it yields the same E/M steps of the EM algo! So, we will look at the EM algo.

# Blank space for illustrations

- Inference probability?  $r_k^{(n)} = p(z^{(n)}=k \mid x=x^{(n)}; \theta^t)$

Recall: MLE for one 1D Gaussian (closed-form soln.) –  
*How to change it to learn mixture of Gaussians?*

- Log likelihood:

$$\ln P\left(\left\{x^{(n)}\right\}_{n=1}^N \mid \mu, \sigma\right) = -N \ln(\sqrt{2\pi}\sigma) - \sum_n (x^{(n)} - \mu)^2 / (2\sigma^2)$$

- MLE estimates:

$$\hat{\mu} = \sum_{n=1}^N \frac{x^{(n)}}{N}, \quad \hat{\sigma}_N^2 = \frac{\sum_{n=1}^N (x^{(n)} - \hat{\mu})^2}{N}$$

Recall: MLE for one 1D Gaussian (closed-form soln.) –  
*How to change it to learn mixture of Gaussians?*

- Log likelihood:

$$\ln P\left(\left\{x^{(n)}\right\}_{n=1}^N \mid \mu, \sigma\right) = -N \ln(\sqrt{2\pi}\sigma) - \sum_n (x^{(n)} - \mu)^2 / (2\sigma^2)$$

$$\theta^t = \underset{\text{current}}{\text{params.}} \cdot r_k^{(n)} = p(z^{(n)}=k \mid x=x^{(n)}; \theta^t)$$

- MLE estimates:

$$\hat{\mu}_N^{(k)} = \frac{\sum_{n=1}^N r_k^{(n)} x^{(n)}}{\sum_{n=1}^N r_k^{(n)}}, \quad \hat{\sigma}_N^{(k)} = \frac{\sum_{n=1}^N r_k^{(n)} (x^{(n)} - \hat{\mu})^2}{\sum_{n=1}^N r_k^{(n)}}$$

(wted. sample mean & variance)

# EM-algorithm in Action: Parameter-learning of a Mixture of 1D Gaussians (aka univariate GMM)

Calculate:

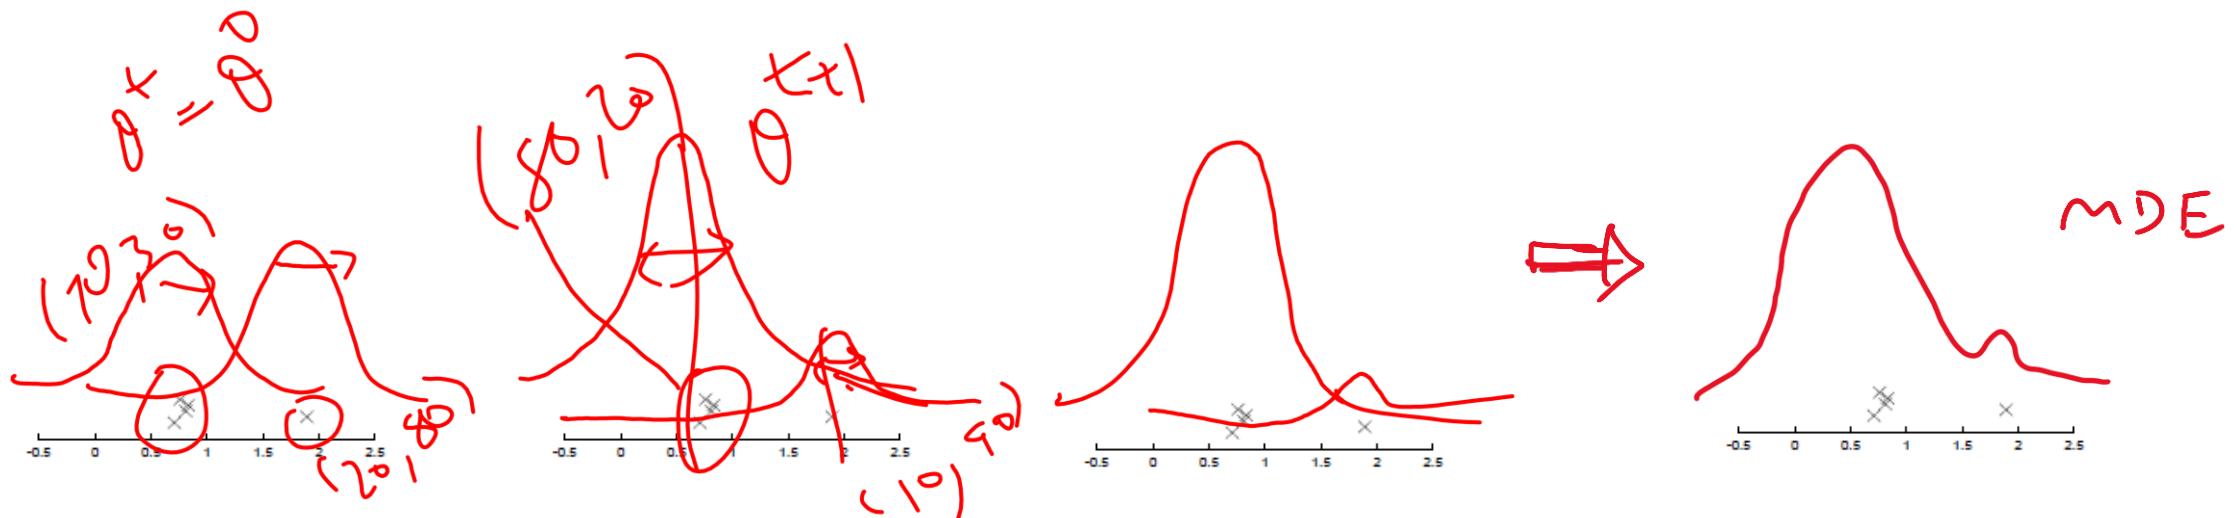
$$r_k(n) = p(z=k \mid \theta_t) = \frac{\pi_k N(x \mid \mu_k, \sigma_k^2)}{\sum_{k=1}^K \pi_k N(x \mid \mu_k, \sigma_k^2)}$$

E-step (Assign)

$$\hat{\mu}_k = \frac{\sum_{n=1}^N r_k(n) x_n}{R_k}; \hat{\sigma}_k^2 = \frac{\sum_{n=1}^N r_k(n) (x_n - \hat{\mu}_k)^2}{R_k}$$

$$\pi_k = \frac{R_k}{\sum_{k=1}^K R_k} = \frac{R_k}{N} \quad (\theta_{t+1})$$

M-step (Update)



Blank space for illustrations

# K-means Clustering Outline

- 1) hard version: algorithm and analysis
- 2) soft version vanilla (aided by intuitive formula)
- 3) soft version vanilla (aided by a vanilla probabilistic model)**
  - 3a) (Brief) Background: Mixture density estimation via EM algorithm
  - 3b) Methodology: Soft k-means heuristic/algorithm**
- 4) soft version enhancements (aided by an enhanced probabilistic model)

# Soft K-means motivation: clustering == mixture density estimation

- Soft K-means: Views clustering problem as inferring hidden cluster labels in a mixture model of two Gaussians.
- Inference task: What is  $P(Z = k | X = x, \theta)$ 
  - **Answers our final question of interest, given that  $\theta$  (its optimal/MLE value) is known**
  - Also, a subroutine for the important task of MLE below!
- Parameter learning: Use MLE using EM algorithm to find the optimal value of  $\theta$ :

$$\gamma_k^{(n)} = P(Z^{(n)} = k | X = x^{(n)}, \theta)$$

$$\hat{\theta}_{ML} = \operatorname{argmax}_{\theta} P(\text{Data} | \text{Model}_{\theta}) = \prod_{n=1}^N (\pi \mathcal{N}(x_n | \mu_1, \sigma_1^2) + (1 - \pi) \mathcal{N}(x_n | \mu_2, \sigma_2^2))$$

Blank space for illustrations

# Simplifying assumptions abt. GMM

Assume:

- 1) MM of 1D Gaussians ( $I=1$ ) ( $x^{(n)} \in \mathbb{R}^I$ )
- 2)  $\sigma_1 = \sigma_2 = \dots = \sigma_K = \sigma$  (known)
- 3)  $\pi_1 = \pi_2 = \dots = \pi_K = \frac{1}{K}$  (known)

# EM-algorithm in Action: Parameter-learning of a Mixture of 1D Gaussians (aka univariate GMM) – BEFORE SIMPLIFICATION

Calculate:

$$r_k(n) = p(z=k \mid \theta_t) = \frac{\pi_k N(x \mid \mu_k, \sigma_k^2)}{\sum_{k=1}^K \pi_k N(x \mid \mu_k, \sigma_k^2)}$$

E-step (Assign)

M-step (Update)

$$\hat{\mu}_k = \frac{\sum_{n=1}^N r_k(n) x_n}{R_k}; \hat{\sigma}_k^2 = \frac{\sum_{n=1}^N r_k(n) (x_n - \hat{\mu}_k)^2}{R_k}$$

$$\pi_k = \frac{R_k}{\sum_{k=1}^K R_k} = \frac{R_k}{N} \quad (\theta_{t+1})$$



# EM-algorithm in Action: Parameter-learning of a Mixture of 1D Gaussians (aka univariate GMM) – AFTER SIMPLIFICATION

Calculate:

$$r_{k(n)} = p(z=k \mid \theta_t) = \frac{N(x \mid \mu_k, \sigma^2)}{\sum_{k=1}^K N(x \mid \mu_k, \sigma^2)}$$

E-step (Assign)

$$\hat{\mu}_k = \frac{\sum_{n=1}^N r_{k(n)} x_n}{R_k}; \quad \text{For all } k=1, \dots, K$$

$$\sigma_k = \sigma(\text{known})$$

$$\pi_k = \frac{1}{K} (\text{known})$$

$$(\theta_{t+1})$$

M-step (Update)

=

$$= \frac{e^{-\beta d(x, \mu_k)}}{\sum_k e^{-\beta d(x, \mu_k)}} \quad \textcircled{1}$$



Blank space for illustrations

# EM algo. – GMM version (aka) Vanilla Soft K-means

**Initialization:** Set  $K$  means  $\{m^{(k)}\}$  to random values.

**Assignment:**  $r_k^{(n)} = \frac{\exp(-\beta d(\mathbf{m}^{(k)}, \mathbf{x}^{(n)}))}{\sum_{k'} \exp(-\beta d(\mathbf{m}^{(k')}, \mathbf{x}^{(n)}))}$ . ① (here  $\beta = 1/(2\sigma^2)$ )  
(E-step)  $m^{(k)} := \mu_k$

**Update:**  $\mathbf{m}^{(k)} = \frac{\sum_n r_k^{(n)} \mathbf{x}^{(n)}}{R^{(k)}}$  ②  
(M-step)

where  $R^{(k)}$  is the total responsibility of mean  $k$ ,

$$R^{(k)} = \sum_n r_k^{(n)}.$$

Details: Same algo. works for multiple dimensions ( $I > 1$ )!

Assuming Gaussians across the  $I$  dimensions are **indept.** for each cluster, i.e.,

$$P(x \mid Z = k', \theta) = \prod_{i=1}^I P(x_i \mid Z = k', \theta), \text{ for any cluster } k'; \text{ then}$$

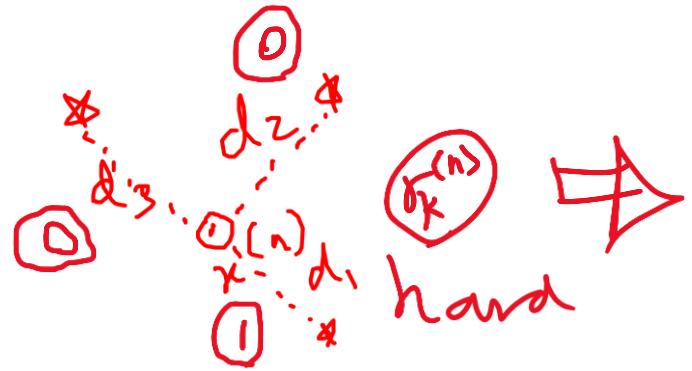
1) Likelihood = ? (1D vs.  $I > 1$ D)

$$P(X = x \mid \theta = \{\pi_1, \dots, \pi_K, m^{(1)}, \dots, m^{(K)}, \sigma\}) = \sum_{k'=1}^K \pi_{k'} \frac{1}{\sigma \sqrt{2\pi}} \exp \left\{ -\frac{(x - m^{(k')})^2}{2\sigma^2} \right\} = \sum_{k'} \pi_{k'} \frac{1}{\sigma \sqrt{2\pi}} \exp(-\beta d(x, m^{(k')}))$$

$$P(X = x \mid \theta) = \sum_{k'=1}^K \pi_{k'} \left( \frac{1}{\sigma \sqrt{2\pi}} \right)^I \exp \left\{ -\frac{\sum_{i=1}^I (x_i - m_i^{(k')})^2}{2\sigma^2} \right\} = \sum_{k'} \pi_{k'} \left( \frac{1}{\sigma \sqrt{2\pi}} \right)^I \exp(-\beta d(x, m^{(k')}))$$

2) Responsibilities = ? (if all  $\pi_{k'}$  are the same)  $r_k^{(n)} = \frac{\exp(-\beta d(\mathbf{m}^{(k)}, \mathbf{x}^{(n)}))}{\sum_{k'} \exp(-\beta d(\mathbf{m}^{(k')}, \mathbf{x}^{(n)}))}.$

Recall: Probabilistic mixture modelling based formula *matches* our intuitive formula for  $r_k^{(n)}$



$$\text{soft} \quad \frac{e^{-\beta d_k}}{e^{-\beta d_1} + e^{-\beta d_2} + e^{-\beta d_3}}$$

Blank space for illustrations

# Soft K-means progress check!

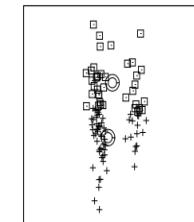
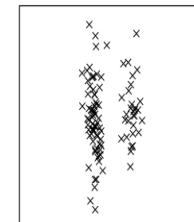
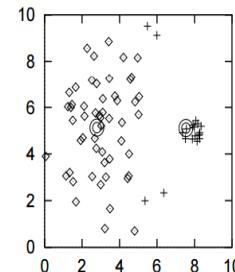
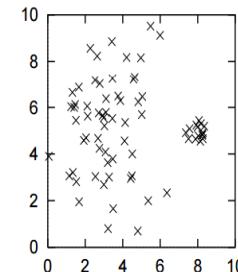
Fixes:

FUZZY ASSIGNMENTS

Doesn't fix:

“Large and small” dataset.

Long and narrow “lozenges” dataset



[DJM]

# K-means Clustering Outline

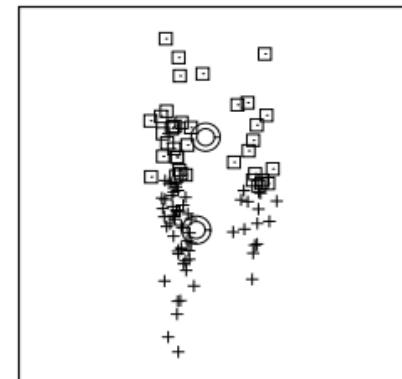
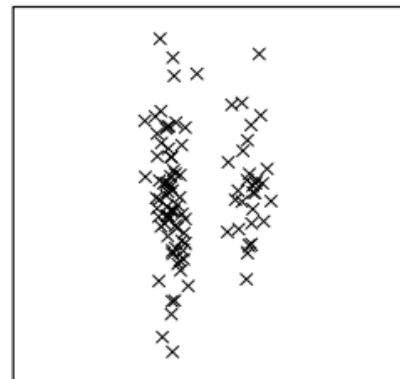
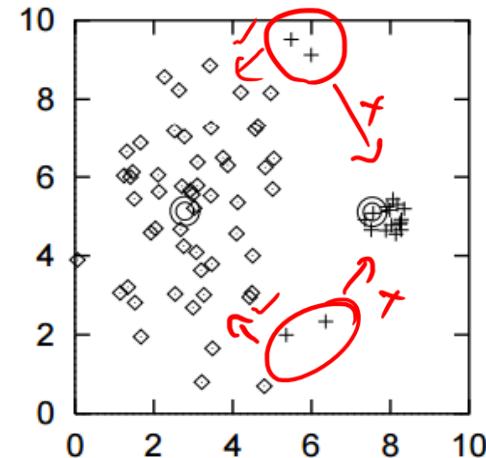
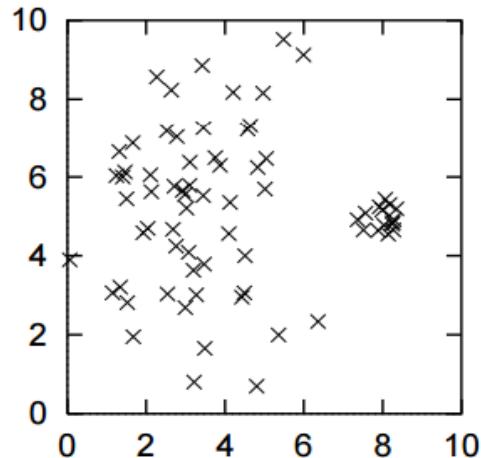
- 1) hard version: algorithm and analysis
- 2) soft version vanilla (aided by intuitive formula)
- 3) soft version vanilla (aided by a vanilla probabilistic model)
  - 3a) Background: Inference+Parameter-learning of a probabilistic mixture model
  - 3b) Methodology: Soft k-means heuristic/algorithm
- 4) soft version enhancements (aided by an enhanced probabilistic model)**

# Enhancing soft kmeans – the idea!

1. Understand the mixture model underlying vanilla soft K-means, and how K-means is simply MLE via EM for this model:
  - i)  $K$  same-width Gaussian  $\sigma$  for each cluster (including same  $\sigma$  across dimensions i.e., spherical Gaussian)
  - ii) Same prior probability for each cluster:  $\pi_1 = \pi_2 = \dots = \pi_K = 1/K$
2. Generalize this model by relaxing these two model assumptions.
  - i) Per-cluster width  $\{\sigma_k\}$
  - ii) Per-cluster prior  $\{\pi_k\}$

Above steps offer a systematic approach to design new variants/enhancements of the K-means algorithm!!!

# Cluster width as extra parameter?



# Soft K-means version 2

(adapt to different  
*cluster widths* and  
*cluster proportions*)

**Assignment:**

$$r_k^{(n)} = \frac{\pi_k \frac{1}{(\sqrt{2\pi}\sigma_k)^I} \exp\left(-\frac{1}{\sigma_k^2} d(\mathbf{m}^{(k)}, \mathbf{x}^{(n)})\right)}{\sum_{k'} \pi_k \frac{1}{(\sqrt{2\pi}\sigma_{k'})^I} \exp\left(-\frac{1}{\sigma_{k'}^2} d(\mathbf{m}^{(k')}, \mathbf{x}^{(n)})\right)}$$

where  $I$  is the dimensionality of  $\mathbf{x}$ .

**Update:**

$$\mathbf{m}^{(k)} = \frac{\sum_n r_k^{(n)} \mathbf{x}^{(n)}}{R^{(k)}}$$

$$\sigma_k^2 = \frac{\sum_n r_k^{(n)} (\mathbf{x}^{(n)} - \mathbf{m}^{(k)})^2}{IR^{(k)}}$$

$$\pi_k = \frac{R^{(k)}}{\sum_k R^{(k)}}$$

where  $R^{(k)}$  is the total responsibility of mean  $k$ ,

$$R^{(k)} = \sum_n r_k^{(n)}.$$

# Soft K-means version 2

(adapt to different  
*cluster widths* and  
*cluster proportions*)

$$\{\sigma_k\}, \{\pi_k\}$$

**Assignment:**

$$r_k^{(n)} = \frac{\pi_k \frac{1}{(\sqrt{2\pi}\sigma_k)^I} \exp\left(-\frac{1}{\sigma_k^2} d(\mathbf{m}^{(k)}, \mathbf{x}^{(n)})\right)}{\sum_{k'} \pi_{k'} \frac{1}{(\sqrt{2\pi}\sigma_{k'})^I} \exp\left(-\frac{1}{\sigma_{k'}^2} d(\mathbf{m}^{(k')}, \mathbf{x}^{(n)})\right)}$$

where  $I$  is the dimensionality of  $\mathbf{x}$ .

**Update:**

$$\mathbf{m}^{(k)} = \frac{\sum_n r_k^{(n)} \mathbf{x}^{(n)}}{R^{(k)}}$$

$$\sigma_k^2 = \frac{\sum_n r_k^{(n)} \underbrace{\|\mathbf{x}^{(n)} - \mathbf{m}^{(k)}\|^2}_{d(x^{(n)}, m^{(k)})}}{IR^{(k)}}$$

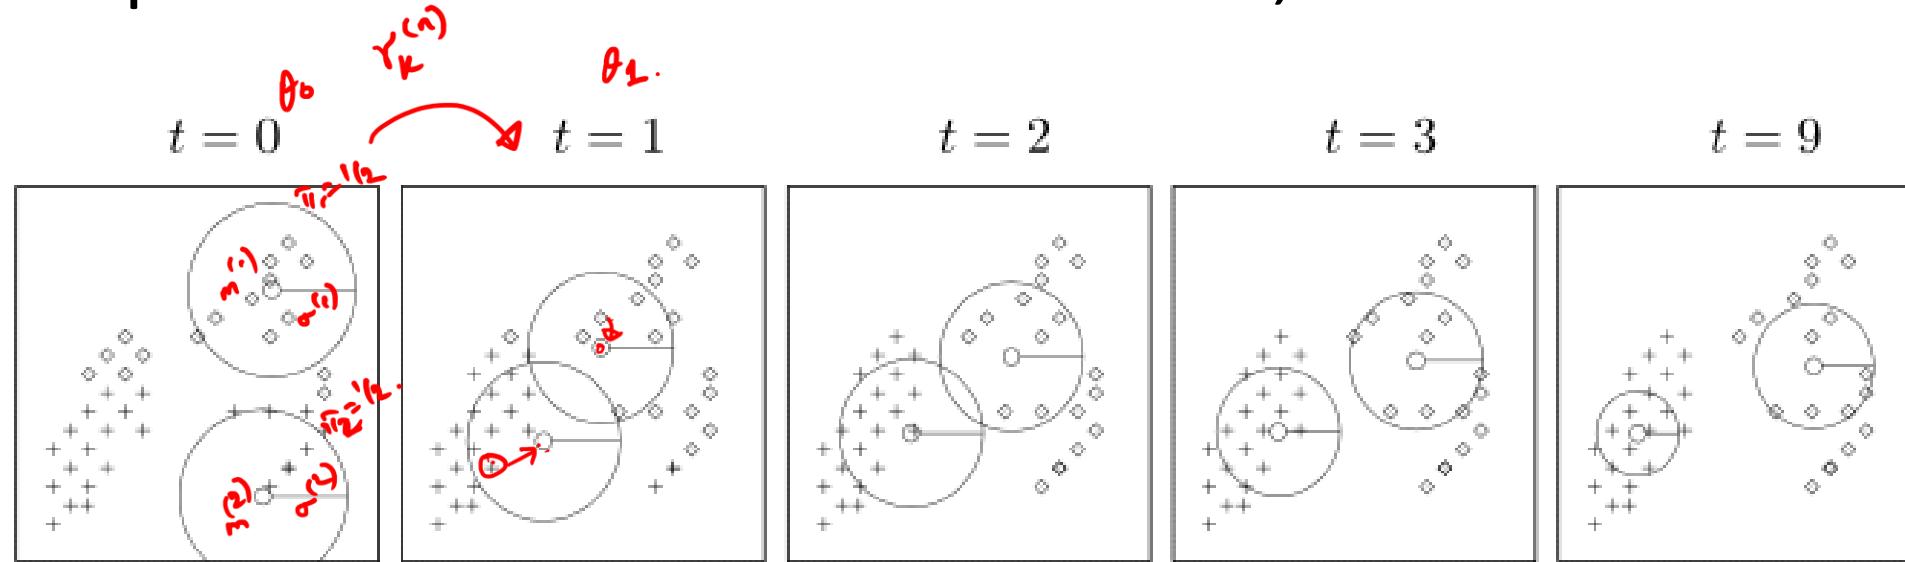
$$\pi_k = \frac{R^{(k)}}{\sum_k R^{(k)}} \approx \frac{r^{(k)}}{N}$$

where  $R^{(k)}$  is the total responsibility of mean  $k$ ,

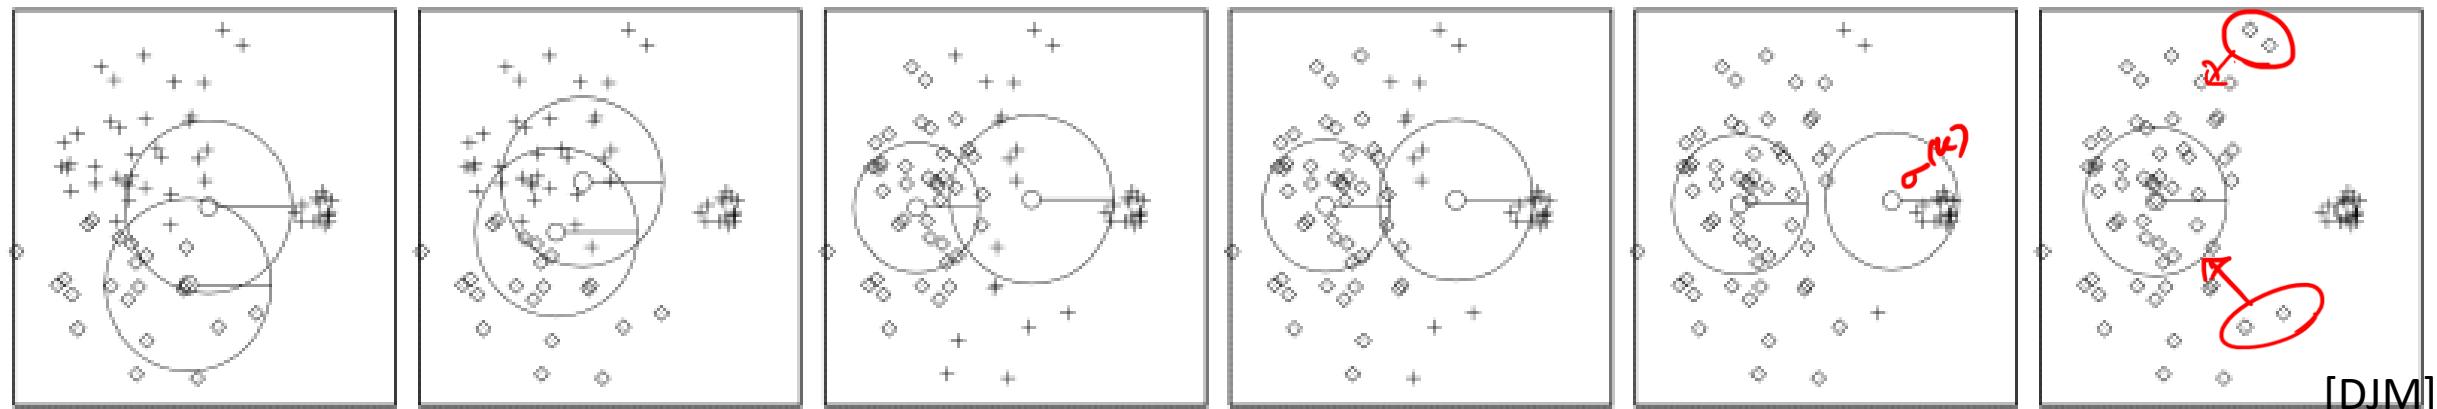
$$R^{(k)} = \sum_n r_k^{(n)}.$$

[DJM]

# Example runs of Soft K-means, version 2



$t = 0$        $t = 1$        $t = 10$        $t = 20$        $t = 30$        $t = 35$

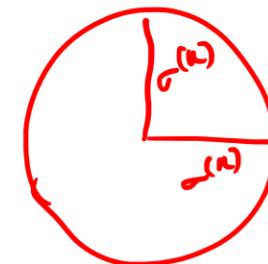


# Final enhancement we will try

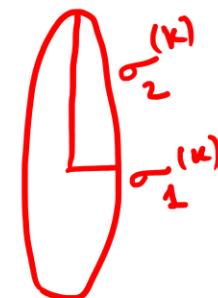
From SPHERICAL to ELLIPSOIDAL axis-aligned (indep.) Gaussians!

So per-cluster, per-dimensionality widths.

$$\sigma \rightarrow \begin{matrix} \sigma^{(k)} \\ \pi_k \end{matrix}$$



each of I dimensions  
also get their own ~~the~~ width.



$$\left\{ \left\{ \sigma_i^{(k)} \right\}_{i=1}^I \right\}_{k=1}^K$$

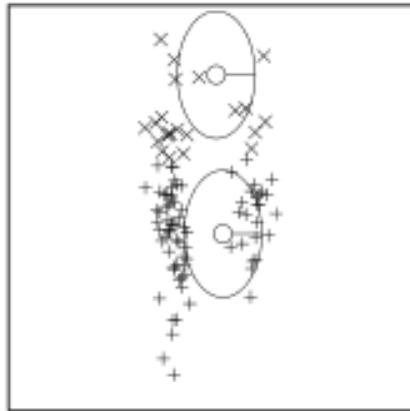
# Soft k-means version 3

$$r_k^{(n)} = \frac{\pi_k \frac{1}{\prod_{i=1}^I \sqrt{2\pi} \sigma_i^{(k)}} \exp \left( - \sum_{i=1}^I (m_i^{(k)} - x_i^{(n)})^2 / 2(\sigma_i^{(k)})^2 \right)}{\sum_{k'} \text{(numerator, with } k' \text{ in place of } k\text{)}} \quad (22.27)$$

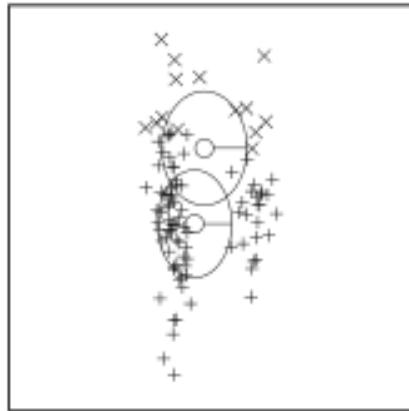
$$\sigma_i^{2(k)} = \frac{\sum r_k^{(n)} (x_i^{(n)} - m_i^{(k)})^2}{R^{(k)}} \quad (22.28)$$

# Example runs of Soft K-means, version 3

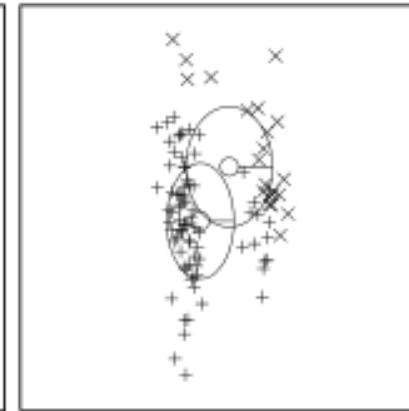
$t = 0$



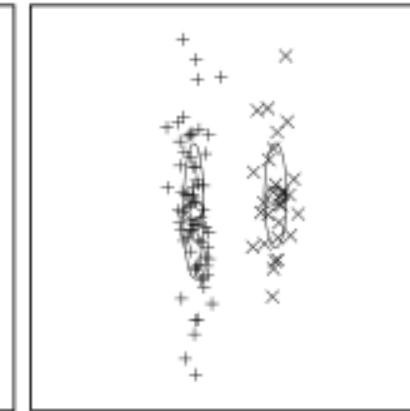
$t = 10$



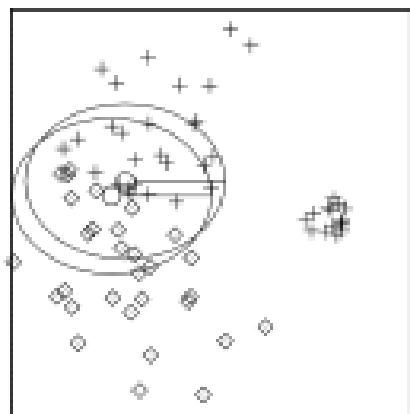
$t = 20$



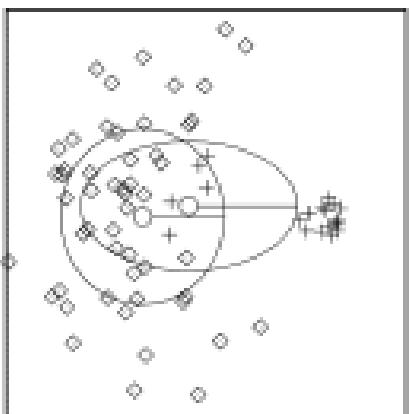
$t = 30$



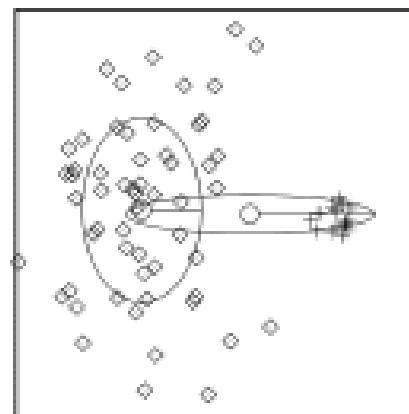
$t = 0$



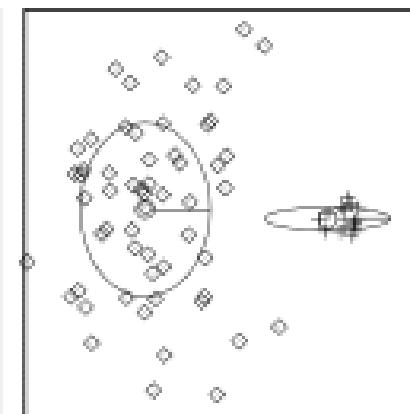
$t = 10$



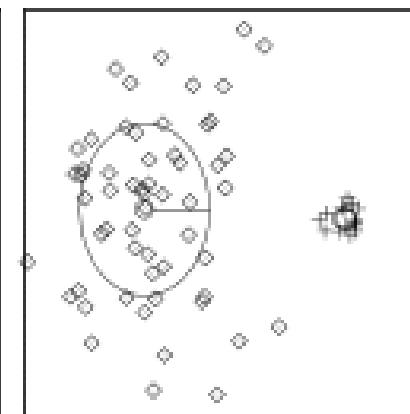
$t = 20$



$t = 26$



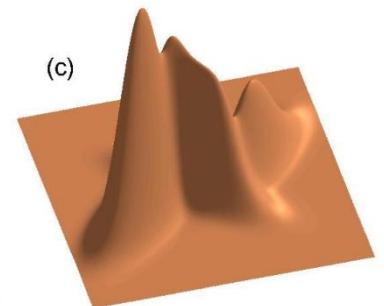
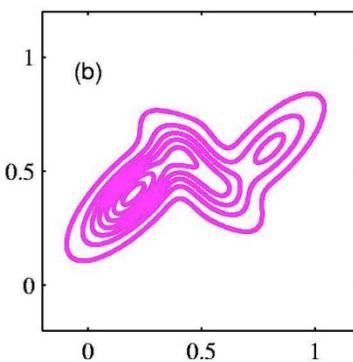
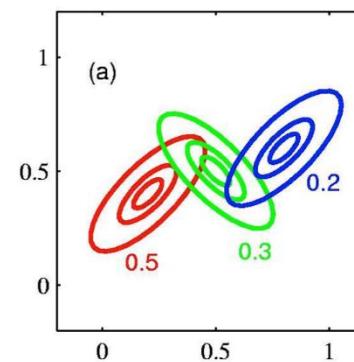
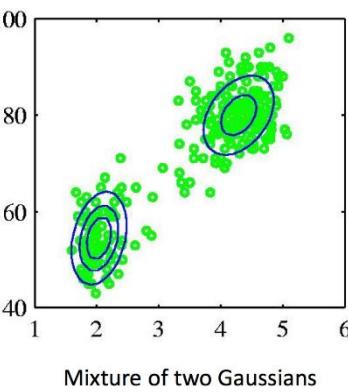
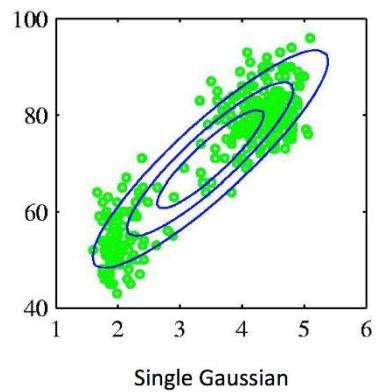
$t = 32$



# Soft k-means version 4?

Q: What if data is not axis-aligned i.e., the different dimensions are not independent?

A: MLE via EM algo. for general GMM



[from ErmongroupStanfordCS228CourseLecNotes]

[CMB]

Example: GMM MLE Learning via EM algo.  
 (aka fully-enhanced Soft k-means algo. or version 4)

**Assignment (E) step:**

$$r_k^{(n)} = \frac{\pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{k'=1}^K \pi_{k'} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_{k'}, \boldsymbol{\Sigma}_{k'})}$$

$\boldsymbol{\mu}_k \in \mathbb{R}^I$   
 $\boldsymbol{\Sigma}_k \in \mathbb{R}^{I \times I}$  (psd)  
 $(\Sigma_k \text{ diag.})$   
 Version 3

**Update (M) step:**

$$\boldsymbol{\mu}_k^{\text{new}} = \frac{1}{R^{(k)}} \sum_{n=1}^N r_k^{(n)} \mathbf{x}^{(n)}$$

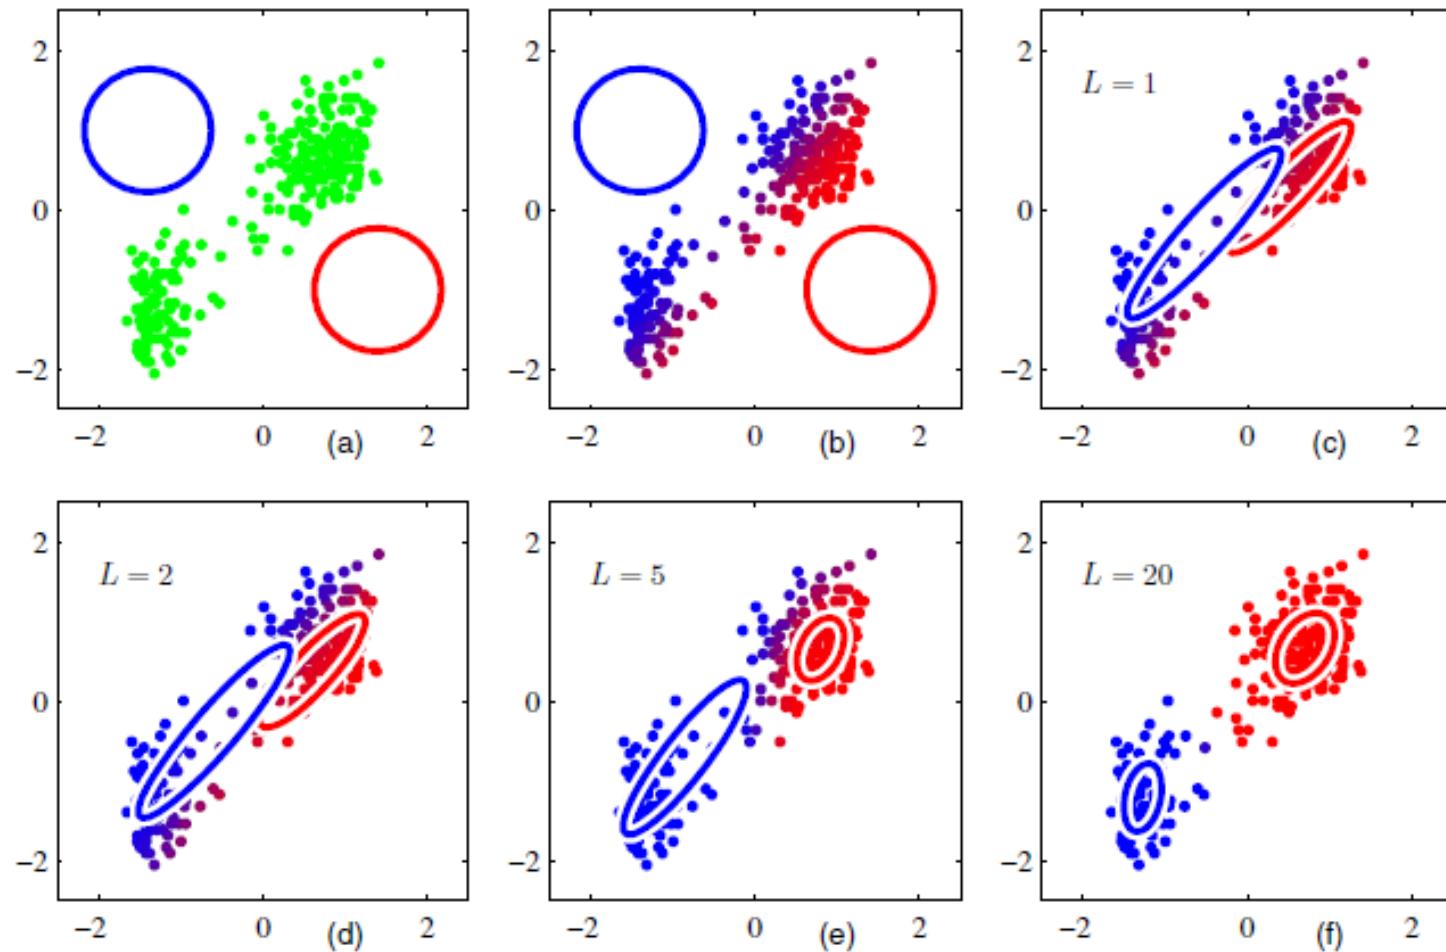
$$\boldsymbol{\Sigma}_k^{\text{new}} = \frac{1}{R^{(k)}} \sum_{n=1}^N r_k^{(n)} (\mathbf{x}^{(n)} - \boldsymbol{\mu}_k^{\text{new}}) (\mathbf{x}^{(n)} - \boldsymbol{\mu}_k^{\text{new}})^T$$

$$\pi_k^{\text{new}} = \frac{R^{(k)}}{N}$$

where

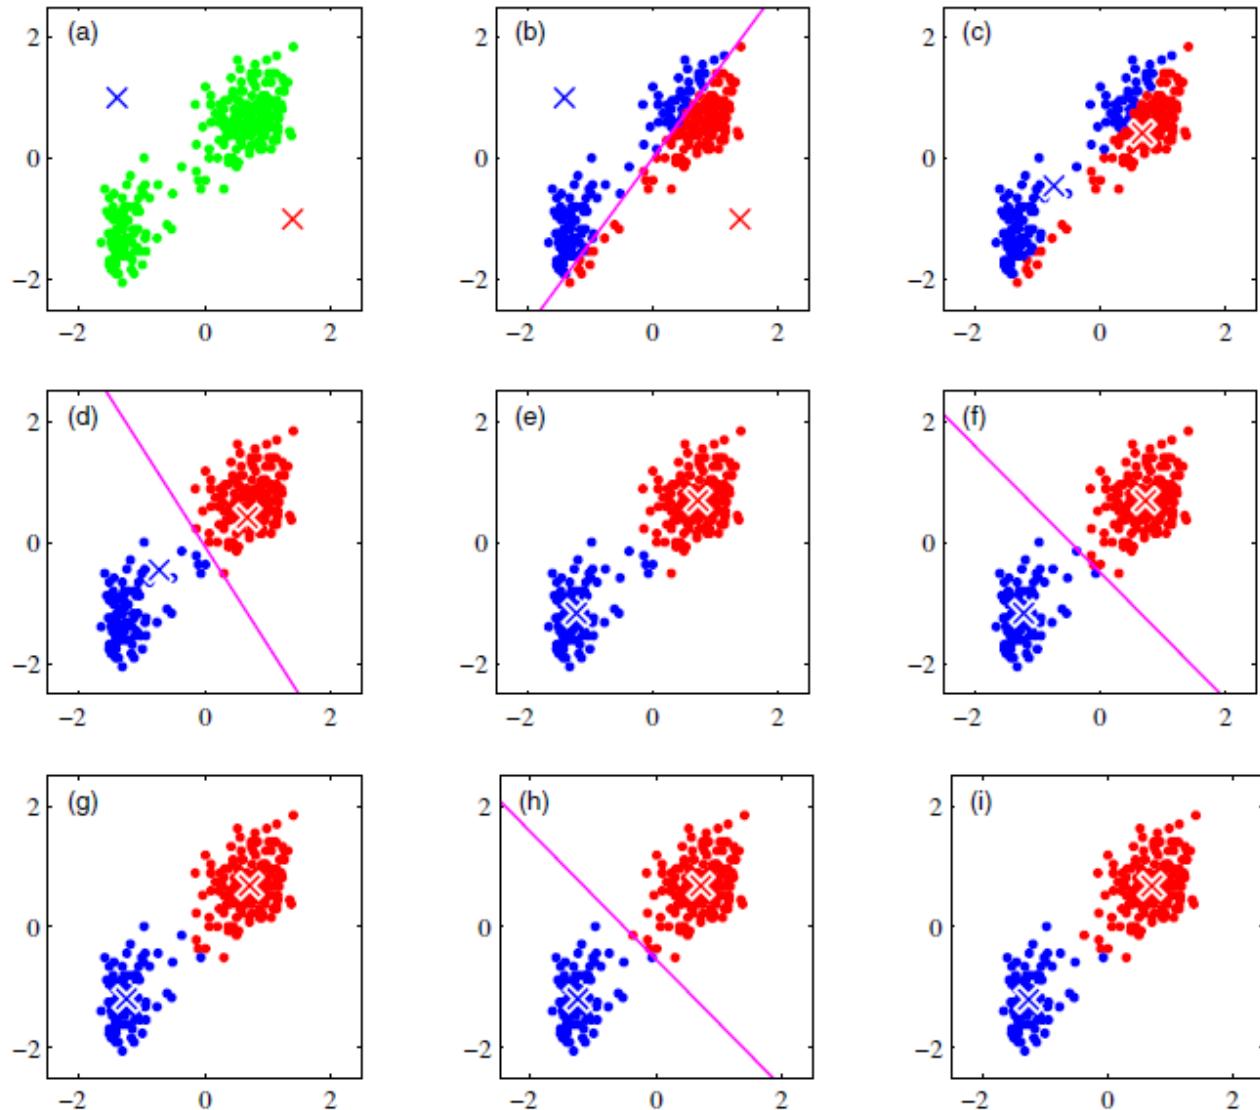
$$R^{(k)} = \sum_{n=1}^N r_k^{(n)}$$

# Fully-enhanced Soft K-means



[CMB]

# Recall: Hard K-means

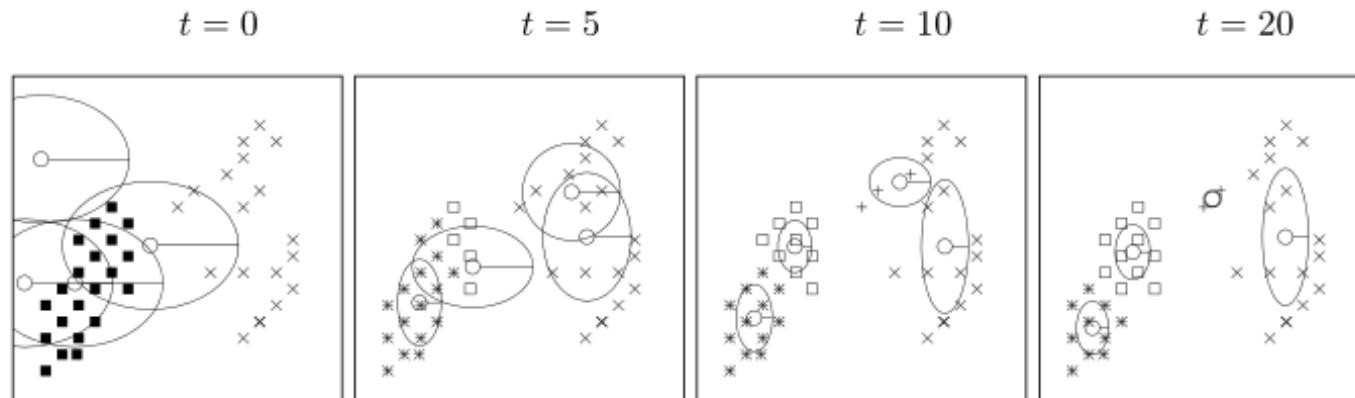


# K-means Clustering Outline

- 1) hard version: algorithm and analysis
- 2) soft version vanilla (aided by intuitive formula)
- 3) soft version vanilla (aided by a vanilla probabilistic model)
- 4) soft version enhancements (aided by an enhanced probabilistic model)

# Summary of “Clustering as Probabilistic Modeling”

- Probabilistic mixture modeling and associated statistical inference gives a systematic approach to design and optimize a data science task.
- Soft K-means successful in many applications (e.g. cancer subtyping), though still not perfect.
  - E.g. MLE blows up when mean sits on exactly one data point. THINK BAYESIAN INFERENCE (aka pseudocounts-like ideas)!



Thank you!